

Introduction to Task-Based Parallel Programming with StarPU

Session 1 — Foundations

Lucas Mello Schnorr, Vinicius Garcia Pinto, Raymond Namyst, Samuel Thibault

CARLA 2026 Advanced Tutorial

<https://github.com/schnorr/starpu-tutorial>

Overview

Where we're going today

- **Session 1 (now)**: the task-based paradigm, and how StarPU implements it — no code yet, just the mental model.
- **Session 2**: your first StarPU programs, from a plain sequential vector scale to task-based, partitioned matrix addition.
- **Session 3**: heterogeneous CPU/GPU codelets, and seeing what actually happened with execution traces (StarVZ).

Setup — start this now, it takes a while

All tutorial material (slides, handouts, exercises, solutions):

`https://github.com/schnorr/starpu-tutorial`

We will use **Podman** throughout (rootless, no background daemon needed — Docker also works, the equivalent command is in the Session 2 material if that is what you have installed):

```
podman pull \
    registry.gitlab.inria.fr/starpu/starpu-docker/starpu:1.4.7
```

This image is **~8 GB** — start the pull **now**, during this lecture, on whatever network you have, so it is already there when Session 2's hands-on begins at 11:00. The exact command to run the container (and what to do if it still refuses to write to your mounted folder) is in the Session 2 material, linked above.

Why task-based programming?

Building a parallel application

Some questions you face when starting to parallelize an application:

- How do I **partition** the work and the data?
- How do the pieces **communicate**?
- How do I **map** pieces onto CPUs/accelerators/nodes?
- If I later add a GPU, or move to a bigger cluster — do I rewrite everything?

The trouble with "classical" parallel programming

With MPI / OpenMP / CUDA written directly:

- You decide, in the code, which core or device runs what, and when.
- Adding a new kind of resource (a GPU, a second node) usually means touching the core logic again.
- Load balancing across heterogeneous, unequal resources becomes your problem, by hand.

A different strategy: describe the work, not the schedule

Task-based programming (also called **dataflow scheduling**):

- The application is split into well-defined **tasks**.
- The program's structure is submitted to a **runtime system** as a **DAG** (Directed Acyclic Graph) of tasks and data dependencies.
- The **runtime** — not your application code — decides scheduling and data movement, and can optimize both.

Advantages and disadvantages

Advantages

- Less parallel-programming complexity exposed to the application.
- The runtime can apply optimizations you didn't have to hand-code.
- Adding a new kind of resource is (mostly) the runtime's problem.

Disadvantages

- One more layer of complexity: the runtime itself.
- You give up some direct control over execution.
- Behavior can be harder to predict (it depends on runtime decisions).

The structure of a task-based application

Three ingredients

Every task-based application, represented as a DAG, is built from:

1. Tasks
2. Data blocks
3. Dependencies

A **unit of computation** — usually sequential code, or a kernel meant to run on an accelerator.

- Operates on well-defined data blocks.
- Can have **multiple implementations** (e.g., one for CPU, one for GPU).

Examples:

- Multiplying two tiles (blocks) of a matrix.
- Applying one step/equation of a simulation model.

A subdivision of the problem's data domain.

- Used as input and/or output of tasks.
- Each block can be read and/or written by different tasks.


Examples:

- Sub-blocks (tiles) of a matrix.
- A subset of elements being simulated.

Dependencies between tasks

Dependencies arise from tasks **sharing the same data**.

- A task of type A produces data consumed by two tasks of type B.



Turning a problem into tasks — is it always easy?

Discretizing an algorithm/problem into tasks and data:

- Similar in spirit to the classic PCAM model (Partitioning, Communication, Agglomeration, Mapping).
- Not every algorithm is a natural fit — but that does not mean an efficient task-based formulation doesn't exist; it may just take more thought to find.

Worked example: finding the largest element

The problem

Given a vector of elements, find the largest one.

What we need to figure out

- What could a **task** be?
- What could the **data** be?
- What **dependencies** result from that choice?

Worked example — the task

Task: find the largest value in a sub-vector.

- Takes a vector of size N .
- Produces: the largest element.
- Call it `MD3` ("max of 3").

Worked example — the data

Data: sub-vectors of the original vector, each of size N .


- For this example, we pick $N = 3$.
- The **outputs** of several MD3 tasks can themselves be grouped into a new sub-vector.

Worked example — the dependencies


Given a vector of 27 elements, we can run $\text{MD } 3$ nine times.

- That produces 9 distinct results.
- Those 9 results can be grouped into 3 new sub-vectors of size 3.
- We can apply $\text{MD } 3$ again on those results — and so on, recursively.

Worked example — the execution graph




Worked example — the DAG



Worked example — where's the parallelism?

Tasks that do **not** share data, and whose dependencies are already satisfied, can run concurrently.



Runtimes

What is a runtime, here?

A **runtime system** executes the DAG you described, instead of you scheduling it by hand.

Widely used in HPC already:

- OpenMP, MPI

Runtimes specifically built around task graphs:

- PaRSEC, OmpSs, OpenMP (tasking, ≥ 4.0), XKaapi, **StarPU**

What do you gain, what do you lose?

Advantages

- Automatic scheduling decisions based on the DAG.
- Automatic mapping of tasks to resources.
- Automatic discovery of parallelism opportunities.
- Automatic data movement between memories.

Disadvantages

- Runtime **overhead**.
- Reduced direct control over execution.


The StarPU runtime

What is StarPU?

A general-purpose, task-based runtime system.

- Targets parallel applications on heterogeneous, multicore machines.
- Developed at Université de Bordeaux / Inria.
- C/C++/Fortran interface for submitting tasks onto resources.

Where StarPU sits in the software stack




StarPU uses the **Sequential Task Flow** (STF) model:

- The application submits tasks **sequentially, as it runs**; the runtime schedules them onto resources **dynamically**.
- The whole DAG does **not** need to be known in advance.
- Tasks can be submitted gradually, as execution proceeds.

Workers

A **worker** is the processing entity that actually executes a task.

- Each computing resource has one or more workers.



Heterogeneous resources need different implementations of the same task.

- Tasks are structured as **codelets**, which may hold one implementation per resource type (CPU, CUDA, OpenCL, ...).
- **Which** implementation actually runs is decided by the runtime, at execution time.

(This is exactly what you will extend by hand in Session 3.)

How StarPU structures a task-based program

- Tasks are described as **codelets**.
- Memory blocks are described as **data handles**.
- Dependencies are computed **implicitly**, from data reuse and submission order — you do not write them explicitly.

(This is exactly what you will build by hand in Session 2.)

Several schedulers are available, selectable at runtime:


- `lws` — local work-stealing
- `eager` — single centralized task queue
- `dm` — deque model, based on the HEFT heuristic
- `dmda` — deque model, data-aware

Selected via the `STARPU_SCHED` environment variable — no code change needed to try a different one.

Analyzing performance

StarPU applications can be analyzed using execution **traces**:

- The FxT format records fine-grained runtime events.
- Traces can be converted to other formats (e.g. Paje) for visualization.



This is exactly what Session 3, Part B, walks you through with StarVZ.

Wrap-up

What's next

You now have the vocabulary for Session 2:

- Codelet, data handle, task, implicit dependency.

In Session 2 you will:

1. Read and run three progressively richer StarPU programs.
2. Write a partitioned block-matrix-addition program from a skeleton.

Questions before the break?